Point Spread Function Estimation and Uncertainty Quantification

Kevin Joyce[†*]

advised by John Bardsley[†] and Aaron Luttman[*]

[†]The University of Montana
[*]National Security Technologies LLC

May 5, 2016

This work was done by National Security Technologies, LLC, under Contract No. DE-AC52-06NA25946 with the U.S. Department of Energy and supported by the Site-Directed Research and Development Program.

Outline

Modeling Imaging Systems

- Convolution with a point spread function
- Estimating the PSF with calibration images

Radial Symmetry for Function Spaces

- Distributions and Sobolev Spaces
- Variable Transformation and the Pullback Operator
- Regularization and Discrete representation

Hierarchical Bayesian Model

- The posterior density
- Gibbs Sampling and Partial Collapse
- Results

Outline

Imaging Model Assumptions

Inverse problem

Ideal Image

Measured image

Image filter

$$b(x) = \iint p(|x-s|) \cdot f(s)\, ds$$

Measurement error

$+\varepsilon \sim$ white noise

Point Spread Function Estimation

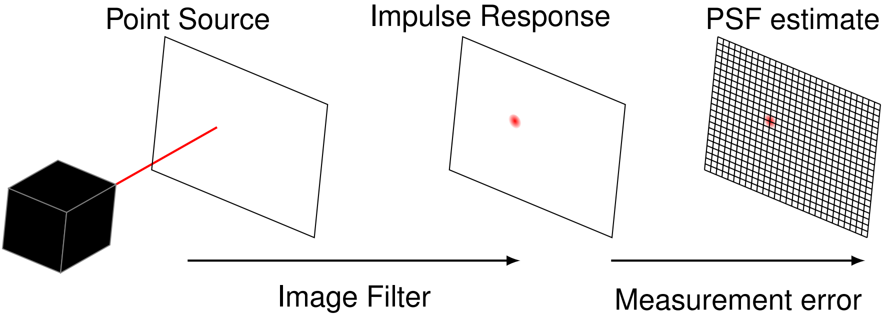

$$\iint p(|x-s|)\delta(s)\,ds$$

$$\iint p(|s|)\delta(x-s)\,ds = p(|x|)$$

$+\varepsilon \sim$ white noise

Point Spread Function Estimation

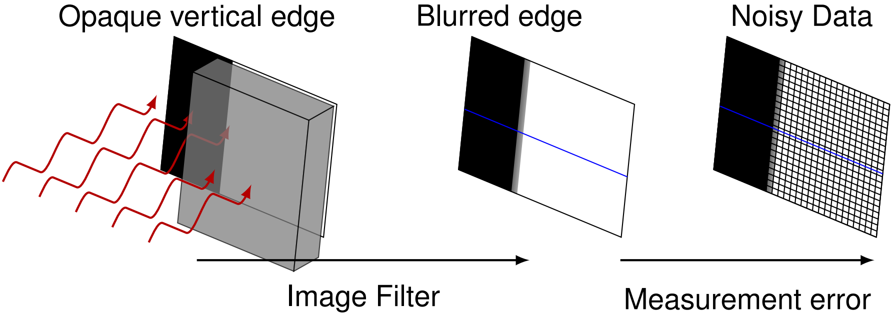

$$\iint p(|x-s|)E(s)\,ds$$

$$\iint p(|s|)E(x-s)\,ds = b(x)$$

$+\varepsilon \sim$ white noise

General Edge Blur Problem

$$b(x,y) = \iint_{\mathbb{R}^2} k(s,t)E(x-s)\,dtds + \varepsilon_{x,y}, \quad E(x) = \begin{cases} 0 & x < 0 \\ 1 & x \geq 0. \end{cases}$$

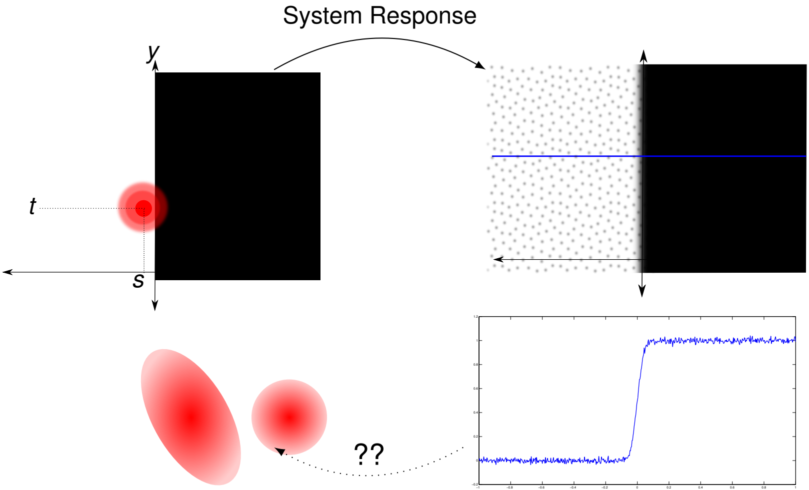

X-ray Edge Calibration Data

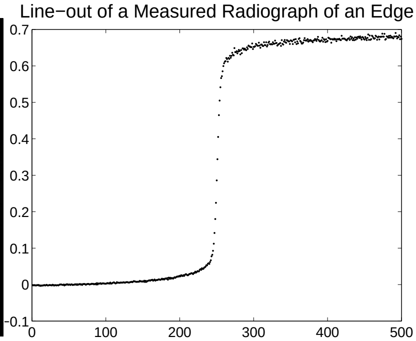

Radiographic data from the Cygnus Dual Beam Radiography Facility at the NNSS in North Las Vegas.

Cygnus X-ray Source

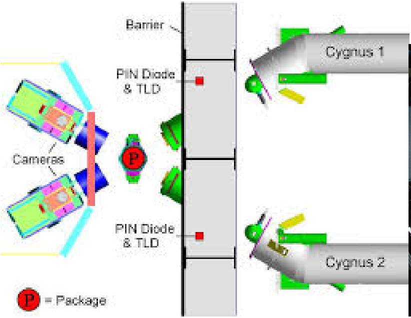

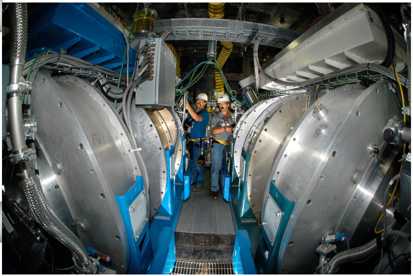

[Smith, et.al. 2012] Cygnus Dual Beam Radiographic Imaging Facility at U1A in the Nevada National Security Site.

Radially Symmetric PSF

We distinguish the radial profile from the kernel by

$k(s,t) = p\left(\sqrt{s^2+t^2}\right)$

$$\begin{aligned} b(x,y) &= \int_{-\infty}^{\infty}\int_{-\infty}^{\infty} k(s,t)E(x-s)dtds + \varepsilon_{x,y} \\ &= \int_{0}^{\infty} p(r)\cdot g(x,r)\,2\pi r dr + \varepsilon_{x,y}. \end{aligned}$$

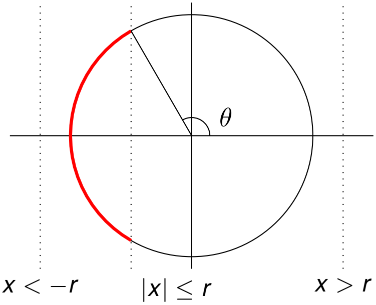

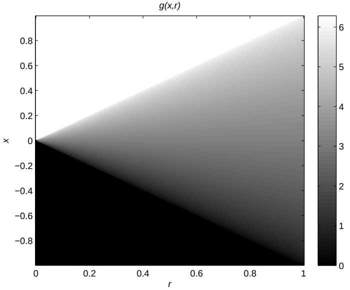

Observe that g is symmetric about $x = 0$.

Outline

Distributions and Sobolev spaces

- Let $\phi \in \mathscr{D}(\Omega)$ denote the space of compactly supported smooth functions defined on an open set $\Omega \subseteq \mathbb{R}^N$, called test functions.
- The space of continuous linear functionals, denoted $f \in \mathscr{D}^*(\Omega)$, are the distributions on Ω, where action of f on ϕ is expressed by $\langle f, \phi \rangle$.
- For functions, the action of the linear functional is $\langle f, g \rangle = \int fg \, dx$.
- Operations are expressed adjointly, e.g. differentiation is given by integration by parts $Df(\phi) \overset{\text{def}}{=} -\langle f, D\phi \rangle$.

Variable Transformation and the Pullback Operator

- Idea: Extend the notion of $k(x, y) = p\left(\sqrt{x^2 + y^2}\right) = T^\sharp p$ to distributions "adjointly" as was done for derivatives

- Use change of variables so that one component is $r = \sqrt{x^2 + y^2}$, for topological reasons, this can only be done on a proper subset of $\mathbb{R}^2$.

- Let $T_{ij}(x, y) = \left(\sqrt{x^2 + y^2}, (-1)^j y\right)$,

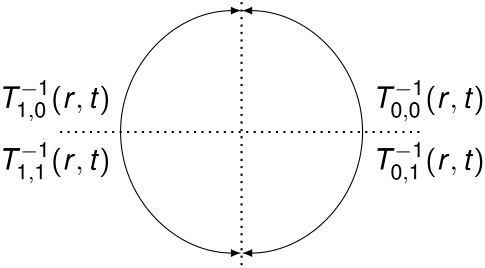

Variable Transformation and the Pullback Operator

- When $k(x,y) = p(\sqrt{x^2+y^2}) = p(T(x,y))$ is a function, then

$$\langle \rho \circ T, \phi \rangle_{\Omega_2} = \sum_{ij} \iint_{Q_{ij}} \rho \circ T(x,y) \cdot \phi(x,y) dxdy$$

$$= \int_0^\infty \rho(r) \left(\int_0^{\sqrt{r}} \sum_{ij} \phi \circ T_{ij}^{-1}(r,t) \left| dT_{ij} \right| dt \right) dr.$$

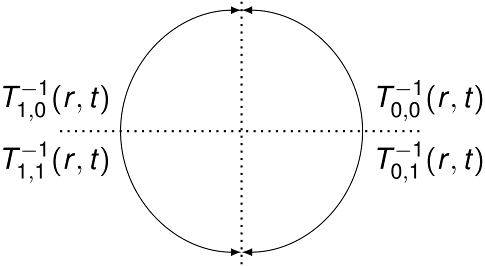

Radial Symmetry for Sobolev Spaces

- The pullback by T on $\mathscr{D}^*(\Omega_1)$ is a linear operator $T^\sharp : \mathscr{D}^*(\Omega_1) \to \mathscr{D}^*(\Omega_2)$ that is injective, continuous, and unique.

- **Definition** $k \in \mathscr{K}^n \subset \mathscr{H}^n(\Omega_2)$, the space of radially symmetric distributions, if there exists a sequence $(\rho_m) \subset \mathscr{D}(\Omega_1)$, so that $(T^\sharp \rho_m)$ is Cauchy with respect to $\|\cdot\|_{\mathscr{H}^k(\Omega_2)}$ and

$$\left\langle k, \phi \right\rangle_{\Omega_2} = \lim_{n\to\infty} \left\langle T^\sharp \rho_n, \phi \right\rangle_{\Omega_2} = \lim_{m\to\infty} \left\langle \rho_m, T_\sharp \phi \right\rangle_{\Omega_1},$$

- **Definition** The space of radial profiles corresponding to $\mathscr{K}^n$ distributions is $\mathscr{P}^n = \{p \in \mathscr{D}^*(\Omega_1) : T^\sharp p \in \mathscr{K}^n\}$.

Radial Symmetry for Sobolev Spaces

- The map $T^\sharp$ induces the inner product

$$(\rho, \omega)_{T(\Omega_1)} = \Big(S_{1/2}(\rho), S_{1/2}(\omega)\Big)_{L^2(\Omega_1)}$$

  where $S(\omega)$ is the shift operator defined by $S(\omega) = \omega(r) \cdot (2\pi r)^{1/2}$.

- When k is a function, the familiar radial transformation is given

$$\iint |k|^2 \, dxdy = \int |p|^2 \, 2\pi r dr.$$

- Moreover, if $\rho, \omega \in \mathscr{D}(\Omega_1)$, then the squared norm of the Laplacian is given by

$$(\nabla T^\sharp \rho, \nabla T^\sharp \omega)_{L^2(\Omega_2)} = (\partial \rho, \partial \omega)_{T(\Omega_1)}.$$

Radial Symmetry for Sobolev Spaces

- So, the expressions

$$(\rho, \omega)_{T(\Omega_1)} = \Big(S_{1/2}(\rho), S_{1/2}(\omega)\Big)_{L^2(\Omega_1)}$$

and

$$(\partial\rho, \partial\omega)_{T(\Omega_1)} = (\nabla T^\sharp \rho, \nabla T^\sharp \omega)_{L^2(\Omega_2)}$$

induce the isometries

$$T^\sharp\Big|_{\mathscr{P}^k} : \mathscr{P}^k \to \mathscr{K}^k \subset \mathscr{H}^k(\Omega_2)$$

.

Tikhonov Laplacian Regularization

For the two representations

$$b = \mathcal{F}k \quad \Longrightarrow \quad b = \left(\mathcal{F}T^{\sharp}\right) p = \mathcal{G}p$$

- Minimizing the second order Tikhonov-Laplacian functional subject to k radially symmetric

$$\frac{\lambda}{2}\left\|b - \mathcal{F}k\right\|_{L^2}^2 + \frac{\delta}{2}\left\langle k, \nabla^4 k\right\rangle_{L^2}$$

- is equivalent to minimizing

$$\frac{\lambda}{2}\left\|b - \mathcal{G}p\right\|_{L^2}^2 + \frac{\delta}{2}\left\langle p, \mathcal{L}^2 p\right\rangle_{T(\Omega_1)}$$

The discrete problem

In order to carry out estimation on a computer, we discretize the integral operator using mid-point quadrature

$$b = \mathcal{G}p + \epsilon \quad \Longrightarrow \quad \boldsymbol{b} = \boldsymbol{G}\boldsymbol{p} + \boldsymbol{\epsilon}$$

Further, we discretize the regularization operator $\mathcal{L}$ using finite differencing

$$\left\langle p, \mathcal{L}^k p \right\rangle_{T(\Omega_1)} = \int \left[\frac{1}{r} \frac{d}{dr} \left(r \frac{d}{dr} \right) \right]^k p(r) r dr$$

$$\Longrightarrow \boldsymbol{L}\boldsymbol{p} = \boldsymbol{r}^{-k+1} \odot \boldsymbol{D} \left(\boldsymbol{r} \odot \boldsymbol{D}\boldsymbol{p} \right)$$

and midpoint quadrature for the inner products

$$\frac{\lambda}{2} \left\| b - \mathcal{G}p \right\|_{L^2}^2 \Longrightarrow \underbrace{\frac{\lambda}{2m}}_{\lambda} \left\| \boldsymbol{b} - \boldsymbol{G}\boldsymbol{p} \right\|_{\mathbb{R}^m}^2$$

and

$$\frac{\delta}{2} \left\| \nabla^2 k \right\|_{L^2} \Longrightarrow \frac{\delta}{2} \left\| \mathcal{L}p \right\|_{rad}^2 \Longrightarrow \underbrace{\frac{\delta}{2n}}_{\delta} \left\| \boldsymbol{L}\boldsymbol{p} \right\|_{\mathbb{R}^n}^2$$

Outline

Hierarchical Model for PSF estimation

Let $\pi(\boldsymbol{x}) = \mathbb{P}(X = x)$ denote the probability density. For λ, δ and $\boldsymbol{p}$ and

$$\boldsymbol{b} = \boldsymbol{G}\boldsymbol{p} + \epsilon$$

assume

- The likelihood $\pi(\boldsymbol{b}|\boldsymbol{p}, \lambda, \delta) = \pi(\boldsymbol{b}|\boldsymbol{p}, \lambda) \propto \lambda^{M/2} \exp\left(-\frac{\lambda}{2}\|\boldsymbol{b} - \boldsymbol{G}\boldsymbol{p}\|^2\right)$ since ϵ is independent Gaussian noise.
- The prior $\pi(\boldsymbol{p}|\delta, \lambda) = \pi(\boldsymbol{p}|\delta) \propto \delta^{N/2} \exp\left(-\frac{\delta}{2}\|\boldsymbol{L}\boldsymbol{p}\|^2\right)$ since $k \sim \mathcal{N}(0, \nabla^{-2}) \implies p \sim N(0, \mathcal{L}^{-2})$
- The hyperpriors $\pi(\lambda) \propto \exp\left(-10^{-4}\lambda\right)$ and $\pi(\delta) \propto \exp\left(-10^{-4}\delta\right)$ are independent Gamma distributions.

Bayesian Posterior

With $\pi(\boldsymbol{b}|\boldsymbol{p},\lambda,\delta), \pi(\boldsymbol{p}|\delta,\lambda)$, and $\pi(\lambda,\delta)$, use Bayes' "Theorem" to obtain

$$\pi_{\mathbf{b}}(\boldsymbol{p},\lambda,\delta) \stackrel{\text{def}}{=} \pi(\boldsymbol{p},\lambda,\delta|\boldsymbol{b}) = \pi(\boldsymbol{b},\boldsymbol{p},\lambda,\delta)/\pi(\boldsymbol{b})$$
$$\propto \lambda^{M/2}\delta^{N/2}\exp\left(-\frac{\lambda}{2}\left\|\boldsymbol{b}-\boldsymbol{G}\boldsymbol{p}\right\|_{\mathbb{R}^m} - \frac{\delta}{2}\left\|\boldsymbol{L}\boldsymbol{p}\right\|_{\mathbb{R}^n} - 10^{-4}(\lambda+\delta)\right)$$

- This is not a "common" probability density, hence simulations from a computer are not readily available.
- Bayes' "Theorem" will allow simulations from the full conditionals $\pi_{\mathbf{b}}(\lambda|\delta,\boldsymbol{p}), \pi_{\mathbf{b}}(\delta|\lambda,\boldsymbol{p})$ and $\pi_{\mathbf{b}}(\boldsymbol{p}|\lambda,\delta)$.
- Because each distribution is from the exponential family, they form a conjugacy such that the full conditionals are "shifts" of the priors.

Full conditional densities

- The resulting expressions are

$$\pi(\lambda|\boldsymbol{b},\boldsymbol{p},\delta) \propto \lambda^{(2N+1)/2+\alpha-1} \exp\left(-\lambda\left(\frac{1}{2}\|\boldsymbol{G}\boldsymbol{x}-\boldsymbol{b}\|^2-\beta\right)\right),$$

$$\pi(\delta|\boldsymbol{b},\boldsymbol{p},\lambda) \propto \delta^{N/2+\alpha-1} \exp\left(-\delta\left(\frac{1}{2}\|\boldsymbol{L}\boldsymbol{p}\|-\beta\right)\right),$$

$$\pi(\boldsymbol{p}|\boldsymbol{b},\lambda,\delta) \propto \exp\left(-\frac{1}{2}\left\langle(\boldsymbol{p}-\boldsymbol{m}_{\lambda,\delta}), \boldsymbol{J}_{\lambda,\delta}(\boldsymbol{p}-\boldsymbol{m}_{\lambda,\delta})\right\rangle\right)$$

where

$$\boldsymbol{J}_{\lambda,\delta} \stackrel{\text{def}}{=} (\lambda\boldsymbol{G}^T\boldsymbol{G}+\delta\boldsymbol{L}) \quad \text{and} \quad \boldsymbol{m}_{\lambda,\delta} \stackrel{\text{def}}{=} \boldsymbol{J}_{\lambda,\delta}^{-1}\lambda\boldsymbol{G}^T\boldsymbol{b},$$

- The matrix solves required for sampling can be efficiently computed using a Cholesky decomposition $\boldsymbol{R}_{\lambda,\delta}^T\boldsymbol{R}_{\lambda,\delta} \stackrel{\text{def}}{=} \boldsymbol{J}_{\lambda,\delta}$ in $O(N^3)$ flops.

Gibbs sampling

The Gibbs sampler [Geman and Geman 1984]:
Given $\left(\lambda^{k-1}, \delta^{k-1}, \boldsymbol{p}^{k-1}\right)$, simulate

1: $\lambda^k \sim \pi_{\mathbf{b}}(\lambda|\delta^{k-1}, \boldsymbol{p}^{k-1})$
2: $\delta^k \sim \pi_{\mathbf{b}}(\delta|\lambda^k, \boldsymbol{p}^{k-1})$
3: $\boldsymbol{p}^k \sim \pi_{\mathbf{b}}(\boldsymbol{p}|\lambda^k, \delta^k)$

The associated Markov Chain is invariant with respect to π:

$$\begin{aligned}
[\mathcal{K}\pi_{\mathbf{b}}](\lambda', \delta', \boldsymbol{p}') &= \iiint \pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda', \delta')\pi_{\mathbf{b}}(\delta'|\lambda', \boldsymbol{p})\pi_{\mathbf{b}}(\lambda'|\delta, \boldsymbol{p})\pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})d\lambda d\delta d\boldsymbol{p} \\
&= \pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda', \delta') \int \pi_{\mathbf{b}}(\delta'|\lambda', \boldsymbol{p}) \int \pi_{\mathbf{b}}(\lambda'|\delta, \boldsymbol{p}) \int \pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})d\lambda d\delta d\boldsymbol{p} \\
&= \mathcal{K}_{\lambda,\delta}\mathcal{K}_{\lambda,\mathbf{b}}\mathcal{K}_{\delta,\mathbf{b}}\pi_{\mathbf{b}}
\end{aligned}$$

$\mathcal{K}_{\delta,\mathbf{b}}\pi_{\mathbf{b}}(\lambda|\delta, \boldsymbol{p}) = \pi_{\mathbf{b}}(\lambda'|\delta, \boldsymbol{p})$ $\quad$ $\mathcal{K}_{\lambda,\mathbf{b}}\pi_{\mathbf{b}}(\delta|\lambda, \boldsymbol{p}) = \pi_{\mathbf{b}}(\delta'|\lambda, \boldsymbol{p})$ $\quad$ $\mathcal{K}_{\lambda,\delta}\pi_{\mathbf{b}}(\boldsymbol{p}|\lambda, \delta) = \pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda, \delta)$

$$\implies \mathcal{K}\pi_{\mathbf{b}} = \pi_{\mathbf{b}}$$

Corollary to Gibbs Invariance

- **Corrollary** Suppose $\mathcal{K} = \mathcal{K}_m \ldots \mathcal{K}_1$ is a transition operator with an invariant density π and $\widetilde{\mathcal{K}}_i$ given $\boldsymbol{x}_{\widehat{i}}$ is an operator such that $\widetilde{\mathcal{K}}_i\big[\pi(x_i|\boldsymbol{x}_{\widehat{i}})\big] = \pi(x_1'|\boldsymbol{x}_{\widehat{i}})$, then $\mathcal{K}_m \ldots \widetilde{\mathcal{K}}_i \mathcal{K}_{i-1} \ldots \mathcal{K}$ is invariant with respect to π.
- **Corrollary** A Metropolis-within-Gibbs algorithm that samples M steps that is invariant with respect to $\pi(x_1'|\boldsymbol{x}_{\widehat{i}})$ is invariant with respect to π.

Gibbs sampling

The Gibbs sampler:
Given $\left(\lambda^{k-1}, \delta^{k-1}, \boldsymbol{p}^{k-1}\right)$, simulate

1. Simulate $\lambda^{k+1} \sim \Gamma\left((2N+1)/2 + \alpha, \frac{1}{2}\|\boldsymbol{G}\boldsymbol{p}^k - \boldsymbol{b}\|^2 + \beta\right)$.

2. Simulate $\delta^{k+1} \sim \Gamma\left(N/2 + \alpha, \frac{1}{2}\left\langle \boldsymbol{p}^k, \boldsymbol{L}\boldsymbol{p}^k \right\rangle + \beta\right)$.

3. Compute $\boldsymbol{R}_{\lambda^{k+1},\delta^{k+1}}, \boldsymbol{m}_{\lambda^{k+1},\delta^{k+1}}$,

   and set $\boldsymbol{p}^{k+1} = \boldsymbol{R}^{-1}_{\lambda^{k+1},\delta^{k+1}}\boldsymbol{z} + \boldsymbol{m}_{\lambda^{k+1},\delta^{k+1}}$ where $\boldsymbol{z} \sim \mathcal{N}\left(\mathbf{0}, \boldsymbol{I}_{N\times N}\right)$.

Correlated δ chains

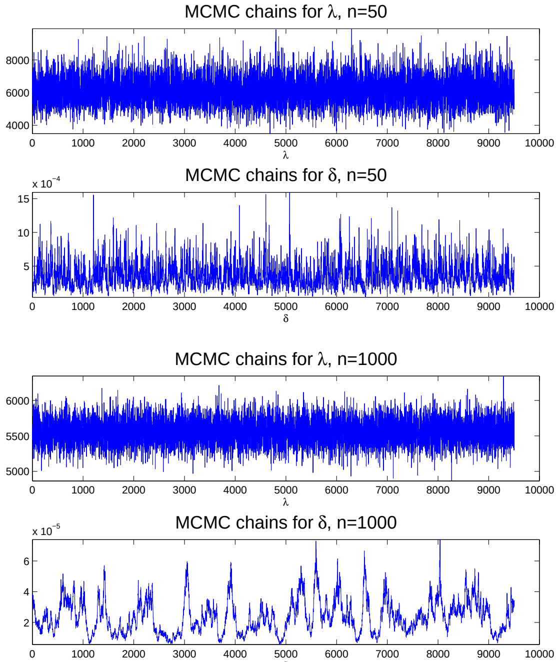

Literature on the Issue

- [Agapiou,Bardsley,Stuart,Papaspiliopoulos, 2014] explained this phenomena theoretically for a general class of Laplacian based Hierarchical samplers for inverse problems.
- The issue arises when the discretization of $\boldsymbol{p}$ closer approximates the continuum, the correlation in the δ component of the Markov Chain becomes more correlated.
- [VanDyke, Park 2008] provide a general method for removing the dependence of problematic components in the Gibbs sampler, called partial collapse.
- The idea has been independently derived in many places, however, if done carelessly [VanDyke, Park 2008] showed that the resulting Markov chain is no longer invariant, although invariance was not proved there.

Marginalized Sampler

Given $\left(\lambda^{k-1}, \delta^{k-1}, \boldsymbol{p}^{k-1}, \widetilde{\boldsymbol{p}}^{k-1}\right)$, simulate

1: $\lambda^k \sim \pi_{\mathbf{b}}(\lambda|\delta^{k-1}, \boldsymbol{p}^{k-1})$
2: $(\delta^k, \widetilde{\boldsymbol{p}}^k) \sim \pi_{\mathbf{b}}(\delta, \widetilde{\boldsymbol{p}}|\lambda^k)$
3: $\boldsymbol{p}^k \sim \pi_{\mathbf{b}}(\boldsymbol{p}|\lambda^k, \delta^k)$

The associated Markov Chain is invariant with respect to $\widetilde{\pi_{\mathbf{b}}}(\lambda, \delta, \boldsymbol{p}, \widetilde{\boldsymbol{p}}) = \pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})\pi_{\mathbf{b}}(\widetilde{\boldsymbol{p}}|\lambda, \delta)$:

$$
\begin{aligned}
[\mathcal{K}\widetilde{\pi_{\mathbf{b}}}] &= \iiiint \pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda', \delta')\pi_{\mathbf{b}}(\delta', \widetilde{\boldsymbol{p}}'|\lambda')\pi_{\mathbf{b}}(\lambda', \boldsymbol{p}|\delta)\widetilde{\pi_{\mathbf{b}}}(\lambda, \delta, \boldsymbol{p}, \widetilde{\boldsymbol{p}})d\lambda d\delta d\boldsymbol{p} d\widetilde{\boldsymbol{p}} \\
&= \pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda', \delta')\pi_{\mathbf{b}}(\delta', \widetilde{\boldsymbol{p}}'|\lambda') \iint \pi_{\mathbf{b}}(\lambda'|\delta, \boldsymbol{p}) \underbrace{\int\int \widetilde{\pi_{\mathbf{b}}}(\lambda, \delta, \boldsymbol{p}, \widetilde{\boldsymbol{p}})d\widetilde{\boldsymbol{p}}d\lambda}_{\pi_{\mathbf{b}}(\delta, \mathbf{p})} d\delta d\boldsymbol{p} \\
&= \underbrace{\pi_{\mathbf{b}}(\boldsymbol{p}'|\lambda', \delta')}_{\pi_{\mathbf{b}}(\mathbf{p}', \delta', \lambda')/\pi_{\mathbf{b}}(\delta', \lambda')} \quad \pi_{\mathbf{b}}(\delta', \widetilde{\boldsymbol{p}}'|\lambda') \underbrace{\iint \pi_{\mathbf{b}}(\lambda', \delta, \boldsymbol{p})d\delta d\boldsymbol{p}}_{\pi_{\mathbf{b}}(\lambda')} \\
&= \pi_{\mathbf{b}}(\boldsymbol{p}', \lambda', \delta')\frac{\pi_{\mathbf{b}}(\widetilde{\boldsymbol{p}}', \lambda', \delta')}{\pi_{\mathbf{b}}(\lambda', \delta')} \\
&= \widetilde{\pi_{\mathbf{b}}}(\lambda', \delta', \boldsymbol{p}', \widetilde{\boldsymbol{p}}')
\end{aligned}
$$

Partially Collapsed Sampler

Given $(\lambda^{k-1}, \delta^{k-1}, \boldsymbol{p}^{k-1})$, simulate

1: $\lambda^k \sim \pi_{\mathbf{b}}(\lambda|\delta^{k-1}, \boldsymbol{p}^{k-1})$
2: $\delta^k \sim \pi_{\mathbf{b}}(\delta|\lambda^k)$
3: $\boldsymbol{p}^k \sim \pi_{\mathbf{b}}(\boldsymbol{p}|\lambda^k, \delta^k)$

The associated Markov Chain is invariant with respect to
$\int \widetilde{\pi_{\mathbf{b}}}(\lambda, \delta, \boldsymbol{p}, \widetilde{\boldsymbol{p}}) = \int \pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})\pi_{\mathbf{b}}(\widetilde{\boldsymbol{p}}|\lambda, \delta) = \pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})$.

- The order of the chain matters in the previous arguments.
- Permuting steps 2 and 3 results in a chain that is no longer invariant with respect to $\pi_{\mathbf{b}}$.
- Cyclically permuting results in a different sampler as well, however, this does not practically effect the overall chain, only the first and last steps, e.g. $(\mathcal{K}_1\mathcal{K}_3\mathcal{K}_2)^N = \mathcal{K}_1(\mathcal{K}_3\mathcal{K}_2\mathcal{K}_1)^{N-1}\mathcal{K}_2\mathcal{K}_3$

Marginalized Posterior Density

In order to sample $\pi_{\mathbf{b}}(\delta|\lambda)$, we complete the square of the quadratic form in $\pi_{\mathbf{b}}(\lambda, \delta, \boldsymbol{p})$ and integrate out λ, this results in

$$\pi_{\mathbf{b}}(\delta|\lambda) \propto \exp\left(\underbrace{(n/2)\delta - \ln|\det \boldsymbol{J}_{\lambda,\delta}| - \frac{\lambda}{2}\langle \boldsymbol{b}, \boldsymbol{H}_{\lambda,\delta}\boldsymbol{b}\rangle - 10^{-4}\delta}_{c(\boldsymbol{R},\lambda,\delta)}\right),$$

where $\boldsymbol{J}_{\lambda,\delta} = \lambda \boldsymbol{G}^T\boldsymbol{G} + \delta \boldsymbol{L}$ and $\boldsymbol{H}_{\lambda,\delta} = \boldsymbol{I} - \lambda \boldsymbol{G}\boldsymbol{J}_{\lambda,\delta}^{-1}\boldsymbol{G}^T$.

- To draw from this density, we embed a Metropolis-Hastings algorithm within the Gibbs sampler.
- **Corrollary** The transition operator $\mathcal{K} = \mathcal{K}_{\lambda,\delta}\widetilde{\mathcal{K}_{MH}}\mathcal{K}_{\delta,\mathbf{b}}$ is invariant with respect to $\pi_{\mathbf{b}}$.
- Both constants can be carried out using a Cholesky factorization $\boldsymbol{R}_{\lambda,\delta}^T\boldsymbol{R}_{\lambda,\delta} = J_{\lambda,\delta}$ in $(O(n^3))$ flops, and will be required for each Metropolis step.

Metropolis-Hastings within Gibbs Sampling

- We use n_{MH} Metropolis-Hastings steps using a Gaussian proposal with variance γ.
- Due to numerical overflow issues, all computations are carried out on the log scale.

1. Simulate $\lambda^{k+1} \sim \Gamma\left((2N+1)/2 + \alpha, \frac{1}{2}\|\boldsymbol{G}\boldsymbol{p}^k - \boldsymbol{b}\|^2 + \beta\right)$.
2. Set $\lambda = \lambda^{k+1}, \delta = \delta^k$ and compute $\boldsymbol{R}_{\lambda,\delta}\ \boldsymbol{m}_{\lambda,\delta},$ then $c(\boldsymbol{R}, \lambda, \delta)$.

   For $j = 1 \ldots n_{mh}$

   i. Simulate $w \sim \mathcal{N}(0,1)$ and set $\delta' = \exp(\gamma w + \delta)$

   ii. Compute $\boldsymbol{R}_{\lambda,\delta'}, \boldsymbol{m}_{\lambda,\delta'},$ then $c(\boldsymbol{R}, \lambda, \delta')$.

   iii. Simulate $u \sim U([0,1])$ and

   if $\ln u > \min\left\{0, c(\boldsymbol{R}, \lambda, \delta') - c(\boldsymbol{R}_{\lambda,\delta}, \boldsymbol{m}_{\lambda,\delta}, \delta)\right\}$

   set $\delta = \delta', \boldsymbol{R}_{\lambda,\delta} = \boldsymbol{R}_{\lambda,\delta'}, \boldsymbol{m}_{\lambda,\delta} = \boldsymbol{m}_{\lambda,\delta},$ and $c(\boldsymbol{R}_{\lambda,\delta}, \boldsymbol{m}_{\lambda,\delta}, \delta) = c(\boldsymbol{R}_{\lambda,\delta'}, \boldsymbol{m}_{\lambda,\delta'}$

   Set $\delta^{k+1} = \delta$
3. Simulate $\boldsymbol{z} \sim \mathcal{N}\left(\boldsymbol{0}, \boldsymbol{I}_{N\times N}\right)$ and set $\boldsymbol{p}^{k+1} = \boldsymbol{R}_{\lambda,\delta}^{-1}\boldsymbol{z} + \boldsymbol{m}_{\lambda,\delta}$.

Marginalized δ chains

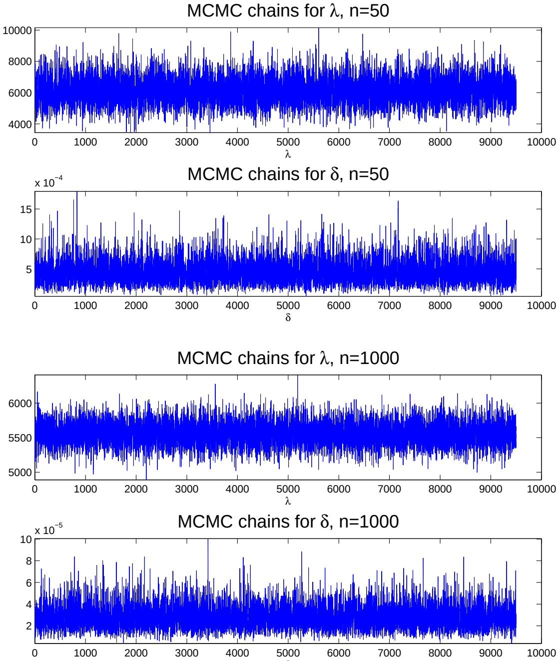

Integrated Autocorrelation

- The variance of the chain-mean estimator for the mean of the invariant density π is

$$\mathrm{Var}(\overline{X}_N) = \frac{1}{N^2}\sum_{k=1}^{N}\mathrm{Var}(X^k) + \frac{1}{N^2}\sum_{k\neq l}^{N}\mathrm{Cov}(X^l, X^k).$$

- When (X_k) are assumed to be nearly identically distributed via the ergodic theorem, then for large N, we can approximate

$$\mathrm{Var}(\overline{X}_N) \approx \frac{\sigma^2}{N^2}\sum_{k=-\infty}^{\infty}\rho(k) \text{ where } \rho(k) \stackrel{\text{def}}{=} \frac{\mathrm{Cov}(X^1, X^{|k|})}{\sigma^2}.$$

- The function ρ is called the normalized auto-correlation function and the parameter $\tau_{\text{int}} \stackrel{\text{def}}{=} \sum_{k=-\infty}^{\infty}\rho(k)$ is the integrated auto-correlation time.

Comparing algorithms

- We can estimate τ_{int} using a convolution estimator.
- The essential sample size is $ESS = M/\widehat{\tau}_{int}$.

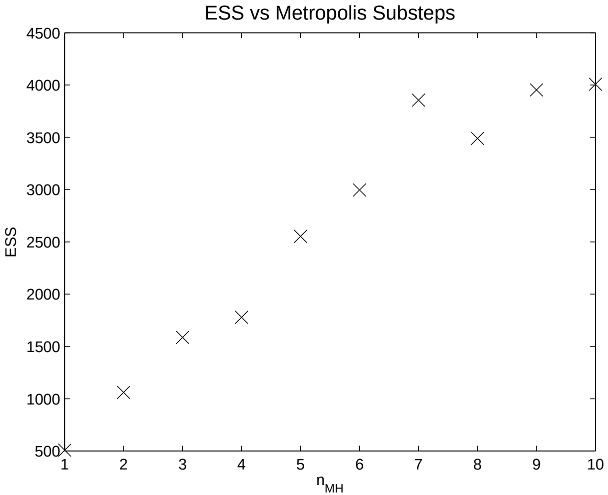

- Our measure of efficiency is $\#Chol/ESS$.

Results

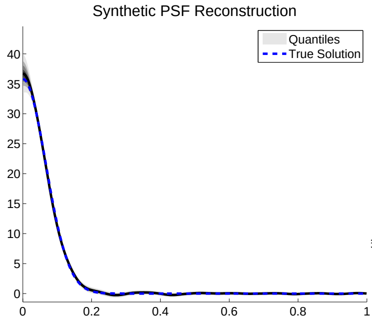

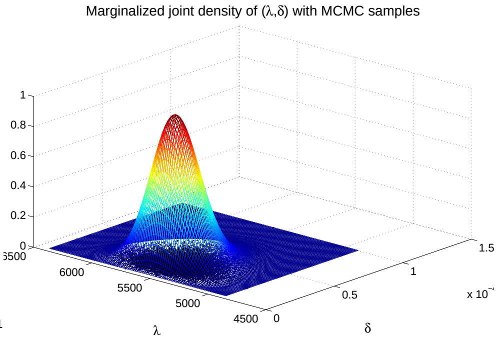

Algorithm	$\widehat{\lambda}_{\mathrm{MCMC}}$ $(\times 10^4)$	$\widehat{\delta}_{\mathrm{MCMC}}$ $(\times 10^{-8})$	λ-p_{Geweke}	δ-p_{Geweke}	IACT	ESS	#Chol/ESS
Gibbs	1.102	6.132	0.998	0.850	36.2	138.0	72.4
PC Gibbs $n_{mh} = 1$	1.102	5.611	0.992	0.943	7.9	633.0	31.6
PC Gibbs $n_{mh} = 5$	1.102	5.515	0.999	0.985	1.3	3799.6	15.8

Results

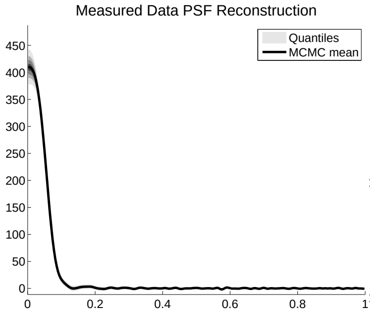

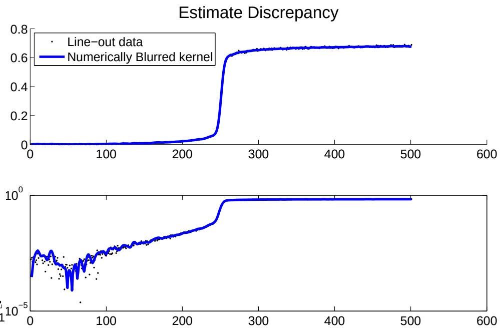

Algorithm	$\widehat{\lambda}_{\text{MCMC}}$ $(\times 10^4)$	$\widehat{\delta}_{\text{MCMC}}$ $(\times 10^{-10})$	λ-p_{Geweke}	δ-p_{Geweke}	IACT	ESS	#Chol/ESS
Gibbs	9.146	1.245	0.995	0.964	14.0	357.6	28.0
PC Gibbs $n_{mh} = 1$	9.167	1.191	0.995	0.998	8.5	587.3	34.1
PC Gibbs $n_{mh} = 5$	9.178	1.189	0.994	0.980	1.5	3278.5	18.3

Summary and Future Work

- We introduced a novel Hierarchical Bayesian non-parametric model for estimating translation invariant and isotropic image blur with and edge.
- We developed the Partially Collapsed Gibbs sampler from the Gibbs sampler, and showed when partial collapse remained stationary.
- We then implemented the algorithm on a synthetic example using Metropolis with Partially Collapsed Gibbs, and showed that it improves the standard Gibbs sampler.
- Future: Develop the model and algorithm completely in infinite dimensions.
- Future: Adapt the strategies to other imaging models that incorporate radial geometry such as Abel and Radon transforms.

Acknowledgements

This work would not have been possible without the support and guidence of my dear friends and colleagues:

- John Bardsley and Aaron Luttman
- Peter Golubtsov, Jon Graham, Leonid Kalachev
- NSTec Cygnus Radiography and Applied Math groups – in particular, Marylesa Howard, Eric Machoro, Tim Meehan, and Steve Mitchell
- Mathetmatical Sciences Faculty and Staff and the Analysis group – Greg St. George, Jen Brooks, Karel Stroethoff, and Linda Azure
- Graduate Chairs – Emily Stone and Cory Palmer
- All fellow graduate students especially my partners in pontification and procrastination Cody Palmer, Nhan Nguyen, and Charlie Katerba
- To all my friends and family without whom I would not be a functioning human – Mom, Dad, Abbey, Dayne, Lora, Alexis, ...